CS 457/657 Database Management Systems

Programming Assignment 1: Metadata Management

Overview

In this assignment, you will write a program that allows a database user to manage the metadata of their relational data. By metadata, we mean the database's high-level information (e.g., database's name, creation time, owner) as well as the properties of the tables (e.g., table's names, attributes, constraints).

System Design

- You are free (in fact, encouraged) to come up with your own design
    - For instance, Sqlite3 uses one single file for each "database."
- Here is one possible design:
    - One Linux directory -> a database
    - One regular file -> a table

Implementation

- The program should not use an external *database* library or an existing SQL parser/compiler.
- Choose your own programming language, e.g., Python, Java, C/C++
    - Please pick one that you are most comfortable/proficient with
    - If you want to choose a language not mentioned above, please contact the TA *before* you start coding
- Functionalities:
    - Database creation, deletion
    - Table creation, deletion, update, and query

Interface

- A similar but simpler interface than Sqlite3
- Examples (on a Linux terminal, `#` denoting the prompt):
    - # ./<your_program> <enter>
      # CREATE DATABASE db_name <enter>;
        - The terminal should output whether the command is successful or failed
            - If failed, don't crash but gracefully report why
        - Then when you check your file system, it might look like this:
            - ~/your_home_directory/cs457/pa1/db_name
    - # ./<your_program> <enter>
      # USE db_1; CREATE TABLE test_tbl (a1 int, a2 char(9));
        - If successful, then your file system might look like this:
            - ~/your_home_directory/cs457/pa1/db1/test_tbl

Testing

- We will test your program on Ubuntu (version 14 or above)
- If your program cannot compile on our testbed, we may ask you to demo your program
  - Try not to use many exotic libraries
- A full test script will be provided, for instance:
  - # ~/cs457/pa1/<your_program> < PA1_test.sql (expect the <u>standard input</u>)
    - Alternatively, you can use a file name as an argument to your program.
    - You will NOT lose points by only supporting a filename-argument interface, but keep in mind that the *standard input* interface would be more desirable for your users (e.g., our TA).
    - If you choose to support only the filename parameter, you will likely need to do more work for a multi-user scenario, which you will support in the 4-th project.
  - You don't need to parse the comment lines (i.e., lines starting with "- -")
  - We will not to test your programs with other scripts/commends
    - However, it's always good to consider more edge cases
    - Try not to hardcode your parser

Grading (20 points total)

- This is an individual assignment.
- Design document that clarifies the followings: (5 points)
  - How your program organizes multiple databases
  - How your program manages multiple tables
  - At a high level, how you implement those required functionalities
- Source code (15 points)
  - Coding style and clarity, 5 points
    - Appropriate parenthesis locations, indention, etc.
    - Always write comments at the beginning of source files
      - Author, date, history, etc.
    - Always write comments at the beginning of non-trivial class/function
      - What this class/function does, high-level algorithm if needed
    - Write in-line comments for non-trivial blocks of code
  - Functionality, 10 points
    - Refer to the test script for detailed breakdowns

Submission

- WebCampus
- Compress all your source code and report into one package in this format:
  - <your_netid>_pa1
- Late penalty: 10% per day